


Step 0: is it easy?



- 1 ▶ **Q:** for “sentiment on brands”?
- 2 ▶ **Q:** for “topics in letters”?
- 3 ▶ **Q:** for “relevance of citations”?

Natural language and ambiguity

- ▶ Text is (usually) natural language
- ▶ Natural means “as humans naturally express” \implies ambiguity!

Expect “raw results” to be worse than in normal ML!

Subsection 1

Sentiment analysis (and text categorization)

Problem formalization

- ▶ Input: a piece of text (*document*)
- ▶ Output?
 - ▶ a numeric value in $[-1, 1]$ (positivity)
 - ▶ a categorical value in $\{\text{Pos}, \text{Neg}\}$
 - ▶ a categorical value in $\{\text{Pos}, \text{Neutral}, \text{Neg}\}$
- ▶ **Q:** learning data?

Regression, multiclass classification, or binary classification (possibly with confidence).

We can use the techniques we already know (e.g., RF)!

Which are the features?

Good coffee. Great for families. Always had good service. We go early so pretty empty. Flexible with menus. Wish they would remove service charge.

Need to transform a document into an numerical vector!

Text to features

- ▶ Not only for sentiment analysis
- ▶ Many options
- ▶ Options can be combined

Bag of words

- ▶ One dimension (feature, dependent variable) for each word
- ▶ Value of $x_{i,j}$ is the number of occurrences of j -th word in the i -th document.

$t_1 = \text{the cat is on the table}$

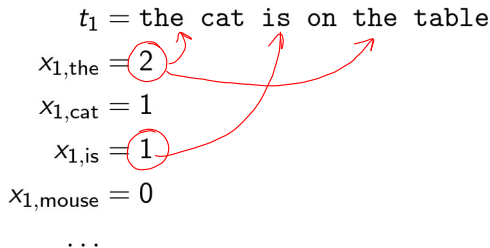
$x_{1,\text{the}} = 2$

$x_{1,\text{cat}} = 1$

$x_{1,\text{is}} = 1$

$x_{1,\text{mouse}} = 0$

...



Bag of which words?

- ▶ One dimension (feature, dependent variable) for each word

Which words? How big is x_i ? $p = ?$

- ▶ common solution: the most k frequent words in the corpus

“Interesting” words not frequent enough in the corpus may be lost

Stop words

- ▶ Some words may be very frequent, but useless for specific task (e.g., sentiment analysis)
 - ▶ a, an, the, are, ... (**stop words**)
- ▶ Just remove them!

Stop words are language dependent!

Stemming

- ▶ There are variants for many words:
 - ▶ drink, drinks, drinking
 - ▶ happy, happier
- ▶ Even more in other languages:
 - ▶ mangio, mangia, mangi, mangiai, ...
- ▶ **Stemming**: reduce word to its word stem (the morphological root)
 - ▶ drinking → drink
 - ▶ argued → argu

Stemming are language dependent!

A typical workflow

- ▶ Preprocessing ($d \rightarrow d'$)
 1. remove punctuation
 2. to lowercase
 3. remove stop words
 4. stemming
- ▶ Learning
 1. preprocess each d in corpus
 2. find most frequent k words in preprocessed corpus
 3. compute \mathbf{X}
 4. learn a classifier
- ▶ Predicting
 1. preprocess input d
 2. predict based on preprocessed d

Limitations and caveat: punctuation

- ▶ remove punctuation

It has been show that often punctuation matter (e.g., Twitter sentiment analysis):

- ▶ I just saw Alice.
- ▶ I just saw Alice!!!!
- ▶ I just saw Alice!!! :-)))))

Case

- ▶ to lowercase

Case may be relevant in some case (e.g., music genre preferences classification):

- ▶ I like the Take That and I hate The Who.
- ▶ Who likes to take that song of Hate? Me!

Goal, context, hypothesis

Twitter profiling: predict age and gender of user from his/her tweets. (Q: what kind of problem/problems?)

- ▶ people of different ages differently use case
- ▶ people of different age/gender differently use punctuation

A step in the workflow corresponds to an (implicit) hypothesis:

- ▶ remove stop words → stop words frequencies is not useful for predicting X

Words that matter

Word count may be too coarse to capture desired information:

- ▶ documents with very different lengths
- ▶ irrelevant terms with general high frequencies

Use frequency or more complex variants

tf-idf

$t \rightarrow \text{term} \rightarrow \text{word}$

$$x_{i,j} = x_{d,t} = \text{tf}(t, d) \text{idf}(t, D)$$

- ▶ $\text{tf}(t, d) = f_{t,d}$, term frequency
 - ▶ the more important the term t in document d , the larger
- ▶ $\text{idf}(t, D) = \log \frac{|D|}{|\{d \in D: f_{t,d} > 0\}|}$, inverse document frequency
 - ▶ the more common t in the corpus, the lower

RARENESS

Bag of words and ordering

Sentiment analysis of restaurant reviews:

t_1 = The beer was good and the pub was not too noisy.

t_2 = The beer was not good and the pub was too noisy.

$x_1 = x_2$

- ▶ fundamental problem: ordering is lost
- ▶ even more fundamental: natural language can be hard to algorithmically understand (irony, sarcasm, ...)

Solutions:

- ▶ ngrams
- ▶ text parsing (NLP)

Aside: collecting data for text classification

Example: irony detection

- ▶ Pavel Savov and Radoslaw Nielek. "Ridiculously Expensive Watches and Surprisingly Many Reviewers: A Study of Irony". In: *Web Intelligence (WI), 2016 IEEE/WIC/ACM International Conference on*. IEEE. 2016, pp. 725–729

ngrams

Instead of counting words, count short (up to n) sequences of words:

- ▶ $x_{1,\text{dog,eat,cat}}$ instead of $x_{1,\text{dog}}$
- ▶ size of data (p) grows dramatically (and is sparser)
- ▶ useful in general for manipulating sequences

Generality of a sentiment classifier

How many sentiment classifier should exist? Words that matter in sentiment should be predefined

- ▶ predefined list of opinion words (positive, negative), i.e., features are those words
- ▶ but context often matter
 - ▶ predictable is good for a car and bad for a movie
 - ▶ features for sentiment analysis in Twitter are likely different than features from sentiment analysis of a early '900 writer's correspondence

There are many pre-trained tool, often with more complex outcome than positive/negative.

Out-of-the-box sentiment analysis

Should I use a pre-trained tool or build my own?

It depends:

- ▶ is sentiment analysis just a piece of a more complex ML system?
- ▶ which is my budget?
- ▶ is learning data easily available?

Parsing: POS tagging

Assign a role to part of speech (**POS**):

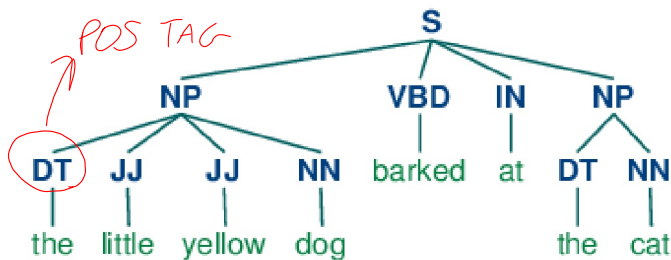


Image from <http://www.nltk.org/>

Lab: Text categorization: sport vs. politics (4h)

Build a binary classifier for tweets: sport vs. politics

1. decide input, output
 2. decide solution assessment
 3. decide (if any) how to obtain learning data
 4. decide workflow and ML technique
- ON PAPER !